
MODULE *Controller*

EXTENDS *Naturals, FiniteSets, Sequences, Messages*

The set of all *ONOS* nodes

CONSTANTS *Nodes*

The current state of *mastership* elections

VARIABLES *term, master, backups*

The current *mastership* event queue for each node

VARIABLE *events*

The current *mastership* state for each node

VARIABLE *mastership*

The unique stream *ID* counter used for correlating controller streams to device streams

VARIABLE *streamId*

Stream change counter used for enforcing state constraints

VARIABLE *streamChanges*

The highest term sent to the device for a node

VARIABLE *sentTerm*

Whether the node has received a *MasterArbitrationUpdate* indicating it is the current master

VARIABLE *isMaster*

Mastership change count used for enforcing state constraints

VARIABLE *mastershipChanges*

A count of all attempted writes to the switch

VARIABLE *writeCount*

Mastership/consensus related variables

$mastershipVars \triangleq \langle term, master, backups, mastershipChanges \rangle$

Mastership arbitration variables

$arbitrationVars \triangleq \langle streamVars, streamId, streamChanges \rangle$

Mastership event variables

$eventVars \triangleq \langle events, mastership \rangle$

Node related variables

$nodeVars \triangleq \langle events, mastership, sentTerm, streamId, streamChanges, isMaster, writeCount \rangle$

This section models the *mastership* election service used by the controller to elect masters. Mastership changes through join and leave steps. Mastership is done through a consensus service, so these steps are atomic. When a node joins or leaves the *mastership* election, events are queued to notify nodes of the *mastership* change. Nodes learn of *mastership* changes independently of the state change in the consensus service.

Returns the set of values in f

$$\text{Range}(f) \triangleq \{f[x] : x \in \text{DOMAIN } f\}$$

Returns a sequences with the element at the given index removed

$$\text{Drop}(q, i) \triangleq \text{SubSeq}(q, 1, i - 1) \circ \text{SubSeq}(q, i + 1, \text{Len}(q))$$

Node 'n' joins the *mastership* election

If the current 'master' is *Nil*, set the master to node 'n', increment the 'term', and send a *mastership* change event to each node. If the current 'master' is non-*Nil*, append node 'n' to the sequence of 'backups'.

$$\begin{aligned} \text{JoinMastershipElection}(n) \triangleq & \\ & \wedge \vee \wedge \text{master} = \text{Nil} \\ & \quad \wedge \text{term}' = \text{term} + 1 \\ & \quad \wedge \text{master}' = n \\ & \quad \wedge \text{backups}' = \langle \rangle \\ & \quad \wedge \text{events}' = [i \in \text{Nodes} \mapsto \text{Append}(\text{events}[i], [\\ & \quad \quad \quad \text{term} \mapsto \text{term}', \\ & \quad \quad \quad \text{master} \mapsto \text{master}', \\ & \quad \quad \quad \text{backups} \mapsto \text{backups}'])]) \\ & \vee \wedge \text{master} \neq \text{Nil} \\ & \quad \wedge \text{master} \neq n \\ & \quad \wedge n \notin \text{Range}(\text{backups}) \\ & \quad \wedge \text{backups}' = \text{Append}(\text{backups}, n) \\ & \quad \wedge \text{events}' = [i \in \text{Nodes} \mapsto \text{Append}(\text{events}[i], [\\ & \quad \quad \quad \text{term} \mapsto \text{term}, \\ & \quad \quad \quad \text{master} \mapsto \text{master}, \\ & \quad \quad \quad \text{backups} \mapsto \text{backups}'])]) \\ & \quad \wedge \text{UNCHANGED } \langle \text{term}, \text{master} \rangle \\ & \quad \wedge \text{mastershipChanges}' = \text{mastershipChanges} + 1 \\ & \quad \wedge \text{UNCHANGED } \langle \text{mastership}, \text{sentTerm}, \text{isMaster}, \text{writeCount}, \text{messageVars}, \text{arbitrationVars} \rangle \end{aligned}$$

Node 'n' leaves the *mastership* election

If node 'n' is the current 'master' and a backup exists, increment the 'term', promote the first backup to master, and send a *mastership* change event to each node. If node 'n' is the current 'master' and no backups exist, set the 'master' to *Nil*. If node 'n' is in the sequence of 'backups', simply remove it.

$$\begin{aligned} \text{LeaveMastershipElection}(n) \triangleq & \\ & \wedge \vee \wedge \text{master} = n \\ & \quad \wedge \vee \wedge \text{Len}(\text{backups}) > 0 \\ & \quad \quad \wedge \text{term}' = \text{term} + 1 \\ & \quad \quad \wedge \text{master}' = \text{backups}[1] \\ & \quad \quad \wedge \text{backups}' = \text{Pop}(\text{backups}) \end{aligned}$$

$$\begin{aligned}
& \wedge \text{events}' = [i \in \text{Nodes} \mapsto \text{Append}(\text{events}[i], [\\
& \qquad \qquad \qquad \text{term} \mapsto \text{term}', \\
& \qquad \qquad \qquad \text{master} \mapsto \text{master}', \\
& \qquad \qquad \qquad \text{backups} \mapsto \text{backups'}])] \\
& \vee \wedge \text{Len}(\text{backups}) = 0 \\
& \quad \wedge \text{master}' = \text{Nil} \\
& \quad \wedge \text{UNCHANGED } \langle \text{term}, \text{backups}, \text{events} \rangle \\
& \vee \wedge n \in \text{Range}(\text{backups}) \\
& \quad \wedge \text{backups}' = \text{Drop}(\text{backups}, \text{CHOOSE } j \in \text{DOMAIN } \text{backups} : \text{backups}[j] = n) \\
& \quad \wedge \text{UNCHANGED } \langle \text{term}, \text{master}, \text{events} \rangle \\
& \wedge \text{mastershipChanges}' = \text{mastershipChanges} + 1 \\
& \wedge \text{UNCHANGED } \langle \text{mastership}, \text{sentTerm}, \text{isMaster}, \text{writeCount}, \text{messageVars}, \text{arbitrationVars} \rangle
\end{aligned}$$

This section models controller-side stream management.

Opens a new stream on the controller side
 $\text{OpenStream}(n) \triangleq$
 $\wedge \text{requestStream}[n].\text{state} = \text{Closed}$
 $\wedge \text{streamId}' = \text{streamId} + 1$
 $\wedge \text{requestStream}' = [\text{requestStream} \text{ EXCEPT } ![n] = [id \mapsto \text{streamId}', \text{state} \mapsto \text{Open}]]$
 $\wedge \text{requests}' = [\text{requests} \text{ EXCEPT } ![n] = \langle \rangle]$
 $\wedge \text{responses}' = [\text{responses} \text{ EXCEPT } ![n] = \langle \rangle]$
 $\wedge \text{streamChanges}' = \text{streamChanges} + 1$
 $\wedge \text{UNCHANGED } \langle \text{mastershipVars}, \text{eventVars}, \text{sentTerm}, \text{isMaster}, \text{responseStream}, \text{writeCount} \rangle$

Closes an open stream on the controller side
 $\text{CloseStream}(n) \triangleq$
 $\wedge \text{requestStream}[n].\text{state} = \text{Open}$
 $\wedge \text{requestStream}' = [\text{requestStream} \text{ EXCEPT } ![n].\text{state} = \text{Closed}]$
 $\wedge \text{sentTerm}' = [\text{sentTerm} \text{ EXCEPT } ![n] = 0]$
 $\wedge \text{isMaster}' = [\text{isMaster} \text{ EXCEPT } ![n] = \text{FALSE}]$
 $\wedge \text{streamChanges}' = \text{streamChanges} + 1$
 $\wedge \text{UNCHANGED } \langle \text{mastershipVars}, \text{eventVars}, \text{responseStream}, \text{messageVars}, \text{streamId}, \text{writeCount} \rangle$

This section models controller-side *mastership* arbitration. The controller nodes receive *mastership* change events from the *mastership* service and send master arbitration requests to the device. Additionally, master nodes can send write requests to the device.

Returns master node 'n' *election_id* for *mastership* term 'm'
 $\text{MasterElectionId}(m) \triangleq m.\text{term} + \text{Cardinality}(\text{Nodes})$

Returns backup node 'n' *election_id* for *mastership* term 'm'
 $\text{BackupElectionId}(n, m) \triangleq m.\text{term} + \text{Cardinality}(\text{Nodes}) - \text{CHOOSE } i \in \text{DOMAIN } m.\text{backups} : m.\text{backups}[i] =$

Returns the *mastership* term for *MasterArbitrationUpdate* 'm'

$$MasterTerm(m) \triangleq m.election_id - Cardinality(Nodes)$$

Node 'n' receives a *mastership* change event from the *mastership* service

When a *mastership* change event is received, the node's local *mastership* state is updated. If the *mastership* term has changed, the node will set a flag to push the *mastership* change to the device in the master arbitration step.

$$\begin{aligned}
LearnMastership(n) &\triangleq \\
&\wedge Len(events[n]) > 0 \\
&\wedge LET \ e \triangleq events[n][1] \\
&\quad m \triangleq mastership[n] \\
&IN \\
&\quad \vee \wedge e.term > m.term \\
&\quad \quad \wedge mastership' = [mastership \text{ EXCEPT } ![n] = [\\
&\quad \quad \quad term \mapsto e.term, \\
&\quad \quad \quad master \mapsto e.master, \\
&\quad \quad \quad backups \mapsto e.backups]] \\
&\quad \vee \wedge e.term = m.term \\
&\quad \quad \wedge mastership' = [mastership \text{ EXCEPT } ![n] = [\\
&\quad \quad \quad term \mapsto e.term, \\
&\quad \quad \quad master \mapsto e.master, \\
&\quad \quad \quad backups \mapsto e.backups]] \\
&\wedge events' = [events \text{ EXCEPT } ![n] = Pop(events[n])] \\
&\wedge UNCHANGED \langle mastershipVars, sentTerm, isMaster, writeCount, messageVars, arbitrationVars \rangle
\end{aligned}$$

Node 'n' sends a *MasterArbitrationUpdate* to the device

If the node has an open stream to the device and a valid *mastership* state, a *MasterArbitrationUpdate* is sent to the device. If the node is a backup, the request's 'election_id' is set to (*mastership* term) + (number of nodes) - (backup index). If the node is the master, the 'election_id' is set to (*mastership* term) + (number of nodes). This is done to avoid *election_ids* ≤ 0 . Note that the actual protocol requires a (*device_id*, *role_id*, *election_id*) tuple, but (*device_id*, *role_id*) have been excluded from this model as we're modelling interaction only within a single (*device_id*, *role_id*) and thus they're irrelevant to correctness. The *mastership* term is sent in *MasterArbitrationUpdate* requests for model checking.

$$\begin{aligned}
SendMasterArbitrationUpdate(n) &\triangleq \\
&\wedge requestStream[n].state = Open \\
&\wedge LET \ m \triangleq mastership[n] \\
&IN \\
&\quad \wedge m.term > 0 \\
&\quad \wedge sentTerm[n] < m.term \\
&\quad \wedge \vee \wedge m.master = n \\
&\quad \quad \wedge SendRequest(n, [\\
&\quad \quad \quad type \mapsto MasterArbitrationUpdate, \\
&\quad \quad \quad election_id \mapsto MasterElectionId(m), \\
&\quad \quad \quad epoch \mapsto m.term]) \\
&\quad \vee \wedge m.master \neq n \\
&\quad \quad \wedge n \in Range(m.backups) \\
&\quad \quad \wedge SendRequest(n, [
\end{aligned}$$

$$\begin{aligned}
& \text{type} \quad \mapsto \text{MasterArbitrationUpdate}, \\
& \text{election_id} \mapsto \text{BackupElectionId}(n, m), \\
& \text{epoch} \quad \mapsto m.\text{term}) \\
& \wedge \text{sentTerm}' = [\text{sentTerm} \text{ EXCEPT } ![n] = m.\text{term}] \\
& \wedge \text{UNCHANGED } \langle \text{mastershipVars}, \text{eventVars}, \text{isMaster}, \text{writeCount}, \text{arbitrationVars}, \text{responses} \rangle
\end{aligned}$$

Node 'n' receives a *MasterArbitrationUpdate* from the device

If the node has an open stream with a *MasterArbitrationUpdate*, determine whether the local node is the master. If the *MasterArbitrationUpdate* 'status' is *Ok*, the 'election_id' matches the last requested *mastership* term, and 'n' is the master for that term, update the node's state to master. Otherwise, the *mastership* request is considered out of date.

Note that the separate 'isMaster' state is maintained to indicate whether the *device* considers this node to be the current master, and this is necessary for the safety of the algorithm. Both the node and the device must agree on the role of the node.

$$\begin{aligned}
& \text{ReceiveMasterArbitrationUpdate}(n) \triangleq \\
& \quad \wedge \text{requestStream}[n].\text{state} = \text{Open} \\
& \quad \wedge \text{HasResponse}(n, \text{MasterArbitrationUpdate}) \\
& \quad \wedge \text{LET } r \triangleq \text{NextResponse}(n) \\
& \quad \quad m \triangleq \text{mastership}[n] \\
& \quad \text{IN} \\
& \quad \quad \vee \wedge r.\text{status} = \text{Ok} \\
& \quad \quad \quad \wedge m.\text{master} = n \\
& \quad \quad \quad \wedge m.\text{term} = \text{MasterTerm}(r) \\
& \quad \quad \quad \wedge \text{sentTerm}[n] = m.\text{term} \\
& \quad \quad \quad \wedge \text{isMaster}' = [\text{isMaster} \text{ EXCEPT } ![n] = \text{TRUE}] \\
& \quad \quad \vee \wedge \vee r.\text{status} \neq \text{Ok} \\
& \quad \quad \quad \vee m.\text{master} \neq n \\
& \quad \quad \quad \vee \text{sentTerm}[n] \neq m.\text{term} \\
& \quad \quad \quad \vee m.\text{term} \neq \text{MasterTerm}(r) \\
& \quad \quad \quad \wedge \text{isMaster}' = [\text{isMaster} \text{ EXCEPT } ![n] = \text{FALSE}] \\
& \quad \wedge \text{DiscardResponse}(n) \\
& \quad \wedge \text{UNCHANGED } \langle \text{mastershipVars}, \text{eventVars}, \text{sentTerm}, \text{arbitrationVars}, \text{requests}, \text{writeCount} \rangle
\end{aligned}$$

Master node 'n' sends a *WriteRequest* to the device

To write to the device, the node must have an open stream, must have received a *mastership* change event from the *mastership* service (stored in 'mastership') indicating it is the master, and must have received a *MasterArbitrationUpdate* from the switch indicating it is the master (stored in 'isMaster') for the same term as was indicated by the *mastership* service. The term is sent with the *WriteRequest* for model checking.

$$\begin{aligned}
& \text{SendWriteRequest}(n) \triangleq \\
& \quad \wedge \text{requestStream}[n].\text{state} = \text{Open} \\
& \quad \wedge \text{LET } m \triangleq \text{mastership}[n] \\
& \quad \text{IN} \\
& \quad \quad \wedge m.\text{term} > 0 \\
& \quad \quad \wedge m.\text{master} = n \\
& \quad \quad \wedge \text{isMaster}[n] \\
& \quad \quad \wedge \text{writeCount}' = \text{writeCount} + 1
\end{aligned}$$

$$\begin{aligned}
& \wedge \text{SendRequest}(n, [\\
& \quad \text{type} \quad \mapsto \text{WriteRequest}, \\
& \quad \text{election_id} \mapsto \text{MasterElectionId}(m), \\
& \quad \text{term} \quad \mapsto m.\text{term}]) \\
& \wedge \text{UNCHANGED } \langle \text{mastershipVars}, \text{eventVars}, \text{arbitrationVars}, \text{isMaster}, \text{sentTerm}, \text{responses} \rangle \\
& \text{Node 'n' receives a write response from the device} \\
& \text{ReceiveWriteResponse}(n) \triangleq \\
& \quad \wedge \text{requestStream}[n].\text{state} = \text{Open} \\
& \quad \wedge \text{HasResponse}(n, \text{WriteResponse}) \\
& \quad \wedge \text{LET } m \triangleq \text{NextResponse}(n) \\
& \quad \text{IN} \\
& \quad \quad \vee m.\text{status} = \text{Ok} \\
& \quad \quad \vee m.\text{status} = \text{PermissionDenied} \\
& \quad \wedge \text{DiscardResponse}(n) \\
& \quad \wedge \text{UNCHANGED } \langle \text{mastershipVars}, \text{nodeVars}, \text{arbitrationVars}, \text{requests} \rangle
\end{aligned}$$

\ * Modification History
\ * Last modified *Thu Feb 21 15:19:30 PST 2019* by *jordanhalterman*
\ * Created *Wed Feb 20 23:49:08 PST 2019* by *jordanhalterman*